

B551 Homework 2

Due date: September 20, 2012

Instructions. Read this assignment carefully, complete the programming assignment and answer all of the written questions. Rename the *username.py* file to *[userID].py*. Place all of your code in the *[userID].py* file and your written answers in the *hw2_answers.txt* file. Submit both files on OnCourse.

Honor Code. You must NOT look at previously published solutions of any of these problems in preparing your answers. You may discuss these problems with other students in the class (in fact, you are encouraged to do so) and/or look into other documents (books, web sites), with the exception of published solutions, without taking any written or electronic notes. Computer code may NOT be shared.

Any intentional transgression of these rules will be considered an honor code violation.

1. Overview

The problems below will ask you to implement three strategies for a game-playing agent for the Gobblet Gobblers game demonstrated in class. You will need to:

- Rename the *players/gobblet/username.py* file to your IU user ID.
- Program your strategies in the indicated functions in that file:
 - An evaluation function giving a minimax value for a game state
 - A minimax search procedure
 - An alpha-beta search procedure
 - A custom search procedure to be used in a game-playing tournament pitting all students' submissions against each other.
- Answer the written questions below.

See *README.txt* for a complete description of the game framework and the role of the file you are modifying.

Submission. Upload your *[userID].py* file and any auxiliary files (additional Python modules or data files you may have used to implement your strategies) to OnCourse. Do not upload modified versions of any of the other provided Python files (in fact, don't modify them at all; when we grade the homework we will use the canonical versions and your changes will not be reflected).

Answers to the written questions should be placed in *hw2_answers.txt* and submitted on OnCourse.

2. The Gobblet Gobblers game

For this homework we are playing a two-player game called "Gobblet Gobblers". It is a variant of tic-tac-toe (or noughts and crosses) with the addition of sized pieces and the ability to move pieces that are already on the board.

Larger pieces can “swallow” and conceal smaller pieces, either the player's own or the opposing player's. The smaller pieces remain on the board, under the larger ones, and are revealed if the larger pieces are moved. When computing the three-in-a-row objective, only the largest piece in a square is considered. The game's complete rules are online and may be viewed at

<http://www.blueorangegames.com/instructions.php>

As the game is capable of cycling (repeating states), we introduce a new rule for our implementation: If a repeated state occurs, the game is declared a draw between the two players. We also adopt the convention that the blue player always moves first and the orange second.

3. Programming

3.1. Setup

Download the file *hw2.zip* from OnCourse. The file contains a directory structure, with base class definitions for a game-playing framework in the files *game.py*, *game_controller.py*, and *game_state.py*. Game definitions for tic-tac-toe and Gobblet Gobblers are in *tictactoe.py* and *gobblet.py*. Some sample game players are defined in the files in the *players* directory.

For more information about these files and their roles, see *README.txt* and the files' documentation. You are principally interested in the *username.py* file in the *players/gobblet* subdirectory, which defines an incomplete Gobblet Gobblers minimax player. You will complete the definition of this player to allow it to play the game.

3.2. Running the program

You may run the program with this command:

```
python game.py gobblet gobblet_simple gobblet_simple
```

Or, on any UNIX/Mac system's command line:

```
./game.py gobblet gobblet_simple gobblet_simple
```

The second and third arguments to the command specify the player modules (Python files) to use for the game; the first specifies the game definition we use. *gobblet_simple.py* resides in the same directory as your Gobblet agent and defines a not-very-smart Gobblet player: It simply selects the first move from the list of successors generated from the current game state. You can also specify *gobblet_human* on the command line as one or both players to allow humans to play the game.

See *README.txt* and the Python files' documentation for more information on running the program, or execute:

```
./game.py -h
```

4. Problems

Question 1 (20 points)

Programming. Implement the evaluation function, `evaluate()`, in your `GobbletPlayer` class.

Given a game state, the function should return a numeric value indicating how “favorable” a game state is to the MAX agent. (A value greater than 0 indicates a state favorable to MAX; less than 0, favorable to MIN; exactly 0, a neutral state with no advantage for either player.)

Notes:

- Consult the definition of the `GobbletState` class (the type of object you'll be evaluating) in `gobblet.py`.
- A popular evaluation system for tic-tac-toe boards is to compare the number of three-in-a-row positions that have not been blocked by the opposing player (and are therefore “available” to the player under consideration). This system is implemented in the `tictactoe_adv.py` player in the `players/tictactoe` directory, which you can feel free to use as a reference.
- However, as pieces can be moved in this game and larger pieces may contain smaller pieces belonging to the other player, the Gobblet Gobblers game is more fluid than the above evaluation suggests. It is strongly suggested to incorporate other features, taking into account available and concealed pieces, in your evaluation.

Written questions. Explain how you devised your state evaluation function, and why you think it is a good measure of a state's favorability to the players.

Question 2 (40 points)

Programming. Implement minimax search in `minimax_move()` in your `GobbletPlayer` class.

Given a current game state, the procedure should search the game tree up to some depth, and use the minimax backup procedure to back up the `evaluate()` function at the leaf nodes back to the root. The return value is a `GobbletMove` object representing the computed move for the current player.

Notes:

- Your search should work whether you are computing the move for MIN or MAX. The current player may be obtained with the provided state's `get_next_player()` function.
- You do not need to write a successor function for this assignment. One is already provided in the `GobbletState.successors()` method. However, when executing the search, the player may only execute a set number of expansions (i.e., you may only call `GobbletState.successors()` a certain number of times per turn before it will stop returning moves.) We use this rather than time to limit players' search depth in the tournament, as time is unreliable and depends on the machine's load among other factors.
- You may obtain the number of available expansions for your current move by calling the provided state's `expansions_count()` function. Obviously, you should stick to calling the provided successor function and not circumvent the limitation on expansions by writing a

separate one of your own. In your minimax function, compute a horizon h and use it to limit the depth of your search so that you don't run out of expansions.

- Again, feel free to consult the Tic-Tac-Toe example player. However, be careful about your math when computing the search horizon: the Tic-Tac-Toe example may not be correct.
- You may want to look into the `mirror()` and `rotations()` functions in *gobblet.py*.

Written questions.

- 1) Given a default limit on the number of expansions, M , what is the largest horizon value h such that minimax is guaranteed to expand a complete tree to that depth (i.e., all branches of the tree would unquestionably go that deep before you ran out of expansions)? Why? Your answer should be an expression involving M .
- 2) Suppose that we had a sufficiently powerful computer to expand the entire Gobblet game tree. What could you conclude from the value of the backed-up minimax value at the root?

Question 3 (40 points)

Programming. Implement the alpha-beta pruning algorithm in `alpha_beta_move()` in your `GobbletPlayer` class. Given a game state, the procedure should execute a minimax search with alpha-beta pruning and return a `GobbletMove` object representing the computed move for the current player. In other words, this function should do all the same things as `minimax_move()`, but with alpha-beta pruning.

Note that again, you must compute a reasonable horizon h , but the horizon may be different than for a simple minimax search.

Written questions.

- 1) Perform some experiments to see what is the largest horizon you can set without encountering the default limit on expansions with your alpha-beta search. (Hint: You may need to use some temporary print statements for this part.) Report on your findings.
- 2) Play your alpha-beta player against your minimax player. What do you observe?
- 3) Suppose computer player A beats computer player B at Gobblet. Would A necessarily win more often against a human player H? Why or why not?

Question 4 (up to 20 points extra credit)

Programming. The search function used during the multi-agent tournament, `tournament_move()`, in your `GobbletPlayer` class is incomplete. Given a game state, the procedure should execute some search routine and return a `GobbletMove` object representing the computed move for the current player.

This function may simply call the minimax or alpha-beta search routines; or you may enhance it with advanced techniques such as singular extension, multiple extension, best-move databases, etc.

Extra credit will be awarded for the agent that wins the tournament by winning the most games. We will also award extra credit for creative implementations of the tournament search.